

# Machine Learning Assignment1

Xiaoxu Gao 4504348

## Question a

The closed form solution of regular ridge regression ( $p=2$ ) is

$$w = (X^T X + \lambda I)^{-1} X^T Y$$

Matlab code:

```
data = load('optdigitsubset.txt');  
dim = 64;  
num = 1125;  
X = data;  
X = [ones(num,1) X];  
Y = [ones(554,1).*(-1);ones(571,1)];  
lambda = 1;  
w = (X'*X+lambda*[zeros(1,dim+1);[zeros(dim,1) eye(dim)]]);)^(-1)*X'*Y;  
w = w(2:end);
```

## Question b

Based on Question a, the following is the Matlab code realted to plot the images.


```
img = reshape(w,[8,8]); %resize the images  
img = mat2gray(img);  
imshow(img,'InitialMagnification','fit'); %fit the screen
```



$\lambda = 10^{-2}$



$\lambda = 10^{-1}$




$\lambda = 1$



$$\lambda = 10^1$$



$$\lambda = 10^2$$




$$\lambda = 10^3$$


### Question c

It's assumed that there are two classes with equal size. The first 554 rows are class0, the rest 544 rows are class1. The following is Matlab code. To better show the tendency, extremely large  $\lambda$  are used here. We can see that for large  $\lambda$ , the solution  $w$  is basically proportional to the difference between the two class means.

```
Matlab code:
data = load('data.txt');
dim = 64;
num = 1108; %every class has 554 samples
X = data((1:num),:);
X = [ones(num/2,1).*(-1);ones(num/2,1)];
lambda = 10^(10);
w = (X'*X+lambda*eye(dim+1))^(10)*X'*Y;
w = w(2:end);
w = mat2gray(w); % the result of w
class0 = X((1:num/2),:);
class1 = X((num/2+1:end),:);
result0 = []; % the means of class0
for i = 1:65
    temp0 = sum(class0(:,i))/(num/2);
    result0 = [result0;temp0];
end
result1 = []; % the means of class1
for j = 1:65
    temp1 = sum(class1(:,j))/(num/2);
    result1 = [result1;temp1];
end
difference = result1 - result0;
difference = difference(2:end) % difference between the two class means
plotmatrix(difference,w); %scatter plot to show the tendency
```



$$\lambda = 10^{10}$$



$$\lambda = 10^7$$

### Question d

$$\min_{w \in \mathbb{R}^d} \sum_{i=1}^N (w^T x_i - y_i)^2 + \lambda \sum_{i=1}^d |w_i|^2$$

The derivative of this formula to  $w_k$  equals:

$$2 \sum_{i=1}^N [x_{ik}(w^T x_i - y_i)] + 2\lambda w_k$$

$$2 \sum_{i=1}^N (w^T x_{ik} - x_{ik}y_i) + 2\lambda w_k$$

Since  $\sum_{j=1}^d x_{ij}w_j$  equals to  $w^T x_i$ , we can get the following formula:

$$2 \sum_{i=1}^N \left[ \left( \sum_{j=1}^d x_{ij}w_j \right) - x_{ik}y_i \right] + 2\lambda w_k$$

### Question e

#### Gradient Descent

step1: start with some  $\theta_0, \theta_1 \dots \theta_{64}$  ( $\theta_0 = \theta_1 = \dots = \theta_{64} = 0$ )

step2: keep changing  $\theta_0, \theta_1 \dots \theta_{64}$  until get minimum value of  $J(\theta_0, \theta_1, \dots \theta_{64})$

Update process:

1.  $temp_0 := \theta_0 - \alpha \frac{d}{d\theta_0} J(\theta_0, \theta_1, \dots \theta_{64})$
2.  $temp_j := \theta_j - \alpha \frac{d}{d\theta_j} J(\theta_0, \theta_1, \dots \theta_{64})$
3.  $\theta_0 = temp_0$
4.  $\theta_j = temp_j$

#### Optimization Toolbox

Matlab toolbox can be used to accomplish this task. Here, I use *fminunc* function to find

the minimum cost.

```
Matlab code:
theta0 = zeros(size(X,2),1);
lambda = 1;
options = optimset('GradObj','on','MaxIter',400);
[theta,J]=fminunc(@(t)(costFunction(t,X,Y,lambda)),theta0,options);


function [cost,grad] = costFunction(theta,X,Y,lambda)
    num = length(Y);
    h = X*theta;
    cost = (1/(2*num))*(sum((h-Y).^2)+lambda*sum(theta(2:end).^2));
    w = zeros(size(theta));
    for i = 1:num
        w = w + (h(i)-Y(i))*X(i,:);
    end
    r = lambda*[0;theta(2:end)];
    grad = (1/num)*(w+r);
end
```

## Question f


Comparison between Normal Equation and Gradient Descent:

	Result	Computation time
<b>Normal Equation</b>	see Questionb	0.157219
<b>Gradient Descent</b>	see Questionf	1.972519

Basically, the difference between these two results are really small. And with the increase of  $\lambda$ , the difference get smaller.



*NormalEquation*( $\lambda = 10^{10}$ )



*GradientDescent*( $\lambda = 10^{10}$ )

## Question g

$$\lambda ||w||_p^p$$

equals:

$$\lambda \sum_{i=1}^d |w_i|^p$$

The derivative of this formula is

$$\lambda p \text{sign}(w_k) |w_k|^{p-1}$$

which equals to  $\lambda p w_k |w_k|^{p-2}$

## Question h

**Extend gradient descent to deal with the general  $L_p$  norm regularizer.**

Based on the formula in Question g, we can get the following code. Since the main program is the same, I only show the code for costFunction.

```
Matlab code:
function [cost, grad] = costFunction(theta, X, Y, lambda)
    num = length(Y);
    h = X*theta;
    cost = (1/(2*num))*(sum((h-Y).^2)+lambda*sum(theta(2:end).^2));
    w = zeros(size(theta));
    for i = 1:num
        w = w + (h(i)-Y(i))*X(i,:);
    end
    p = 3;
    r = lambda*(p/2)*[0; theta(2:end)].*[0; theta(2:end)].^(p-2); %new
    regularization
    grad = (1/num)*(w+r);
end
```

## Question i

In this question, I choose  $\lambda = 10^6, p = 2, p = 1$ . There are two classes, each class has ten samples. Since we are talking about  $p = 2$  and  $p = 1$  here, we need to use the following formular:

$$\lambda p \text{sign}(w_k) |w_k|^{p-1}$$


```

Matlab code:
data = load('data.txt');
dim = 64;
num = 20;
X1 = data(1:20,:);
X2 = data(571:590,:);
X = [X1;X2];
X = [ones(40,1) X];
Y = [ones(20,1).*(-1);ones(20,1)];
lambda = 10^6;
theta0 = zeros(size(X,2),1);
options = optimset('GradObj','on','MaxIter',400);
[theta, J] = fminunc(@(t)(costFunction(t,X,Y,lambda)),theta0,options);
theta = theta(2:end);
plot(theta);


function [cost,grad] = costFunction(theta,X,Y,lambda)
    num = length(Y);
    h = X*theta;
    cost = (1/(2*num))*(sum((h-Y).^2)+lambda*sum(theta(2:end).^2));
    w = zeros(size(theta));
    for i = 1:num
        w = w + (h(i)-Y(i))*X(i,:);
    end
    p = 2; %change the value of p
    r = lambda*(p/2)*sign([0;theta(2:end)]).*[0;theta(2:end)].^(p-1);
    grad = (1/num)*(w+r);
end

```

**Images of solution  $w(\lambda = 10^6)$ :**



$p = 1$



$p = 2$

**When  $p$  gets closer to 1:**

When  $p$  gets closer to 1, the value of cost function gets bigger.

**The reason why some pixel values of  $w$  get closer to zero:**

Here we use regularization to avoid overfitting, so making some  $w$ s small is one of the solutions.

**Question j**

When there are many features but only a few samples, it's highly possible that we will face the problem of overfitting.